

Akash Kondimoena

+1 331-328-8484 | [Email](#) | [Linkedin](#)

Professional Summary

AI/ML Engineer with 4+ years of experience designing, developing, and deploying scalable Artificial Intelligence, Machine Learning, and Generative AI solutions for enterprise applications. Proficient in Python, SQL, TensorFlow, PyTorch, Scikit-learn, and major ML libraries, with expertise in supervised and unsupervised learning, deep learning, neural networks, feature engineering, model training, evaluation, hyperparameter optimization, and end-to-end machine learning pipelines. Experienced in developing NLP, Deep Learning, and Generative AI applications using Large Language Models (LLMs), Prompt Engineering, Retrieval-Augmented Generation (RAG), LangChain, Hugging Face Transformers, OpenAI, and Azure OpenAI to deliver intelligent automation, conversational AI, enterprise AI solutions, and AI-powered decision-making systems, and enterprise knowledge retrieval systems. Skilled in evaluating AI model architectures, implementing Responsible AI and Explainable AI (XAI) practices, and optimizing model performance through continuous monitoring and bias detection. Hands-on experience with Docker, Kubernetes, FastAPI, MLflow, Apache Airflow, CI/CD pipelines, and cloud platforms including AWS, Microsoft Azure, and Google Cloud Platform (GCP). Proven ability to rapidly prototype AI/ML solutions, translate customer and business requirements into technical solutions, assess integration opportunities, troubleshoot technical challenges, and collaborate with cross-functional teams to deliver measurable business impact.

Technical Skills

Programming Languages:

Python, SQL, PySpark

Machine Learning:

Scikit-learn, TensorFlow, PyTorch, Keras, XGBoost, LightGBM, Regression, Classification, Clustering, Recommendation Systems, Feature Engineering, Statistical Modeling, Experimental Design, Hypothesis Testing, Model Training, Model Evaluation, Hyperparameter Tuning

Generative AI & LLMs:

OpenAI GPT, LLaMA, Claude, Gemini, Prompt Engineering, Retrieval-Augmented Generation (RAG), LLM Integration, AI Agents, Agentic Workflows, Tool Calling, Agent Orchestration, Context Management, Memory Systems, Model Context Protocol (MCP), Agent Evaluation Transformers, BERT, Hugging Face Transformers, Text Classification, Named Entity Recognition (NER), Sentiment Analysis, Semantic Search, Text Summarization

Natural Language Processing (NLP):

AI Frameworks & Vector Databases:

LangChain, Hugging Face, Pinecone, FAISS, ChromaDB

APIs & Enterprise Integrations:

APIs & Enterprise Integrations:

Cloud Platforms & MLOps:

AWS (S3, EC2, SageMaker, Glue, Lambda, Redshift), Azure Machine Learning, Google Vertex AI, Docker, Kubernetes, MLflow, Apache Airflow, GitHub Actions, CI/CD

Responsible AI & Model Optimization:

Model Monitoring, Explainable AI (XAI), Bias Detection, Model Performance Optimization

Professional Experience

AI/ML Engineer | Meijer | Chicago, IL, USA

Sep 2024 - Present

- Design and deploy scalable AI/ML solutions on AWS to support demand forecasting, customer behavior analytics, inventory optimization, personalized recommendations, and intelligent retail operations across enterprise applications.
- Develop scalable data pipelines using Amazon S3, AWS Glue, AWS Lambda, and Amazon Redshift to gather, clean, prepare, and perform feature engineering on structured and unstructured datasets for training and evaluating machine learning and Generative AI models.
- Build Generative AI applications using Large Language Models (LLMs), Prompt Engineering, Retrieval-Augmented Generation (RAG), LangChain, OpenAI, and Hugging Face Transformers to deliver conversational AI, semantic retrieval, AI-powered customer support, and intelligent enterprise search solutions.

- Develop Natural Language Processing (NLP) models for sentiment analysis, text classification, document summarization, semantic search, and customer feedback analysis to enhance customer engagement and business decision-making.
- Designed and optimized supervised, unsupervised, and recommendation machine learning models through rapid experimentation, feature engineering, and hyperparameter tuning to improve customer engagement, prediction accuracy, and business outcomes.
- Build, integrate, troubleshoot, and deploy production-ready AI models and REST APIs using Amazon SageMaker, Docker, Kubernetes, and FastAPI; diagnose integration and production issues and resolve technical challenges to maintain scalable, secure, and reliable AI services.
- Integrate enterprise data sources, APIs, embedding models, and vector databases including Pinecone and FAISS to develop AI-powered knowledge management and RAG solutions, identifying integration points across enterprise workflows and customer systems.
- Implement MLOps best practices using MLflow, Apache Airflow, GitHub Actions, and CI/CD pipelines to automate model training, deployment, model monitoring, versioning, lifecycle management, and continuous performance optimization.
- Applied statistical modeling, hypothesis testing, experimental design, and comprehensive model evaluation using Accuracy, Precision, Recall, F1-Score, ROC-AUC, and model drift metrics to improve model performance, reliability, and measurable business impact.
- Collaborate with engineers, AI researchers, product managers, and cross-functional teams to translate business requirements into technical specifications and deliver scalable AI solutions for complex business problems.

Environment: Python, SQL, TensorFlow, PyTorch, Scikit-learn, Pandas, NumPy, Hugging Face Transformers, LangChain, OpenAI, Azure OpenAI, FastAPI, Pinecone, FAISS, AWS (S3, EC2, SageMaker, Glue, Lambda, Redshift), Docker, Kubernetes, MLflow, Apache Airflow, GitHub Actions, Tableau, Power BI.

Data scientist | Enter-Pi Software Solutions | Hyderabad

Sep 2020 – Dec 2023

- Designed and developed end-to-end machine learning solutions for enterprise clients using Python, SQL, Scikit-learn, TensorFlow, and PyTorch to solve business problems through predictive analytics and intelligent automation.
- Built scalable data pipelines and ETL workflows using AWS Glue, AWS Lambda, Python, and SQL to ingest, cleanse, transform, and process structured and unstructured data from multiple enterprise applications.
- Performed exploratory data analysis (EDA) on structured and unstructured datasets to uncover behavioral patterns, identify automation opportunities, engineer predictive features, and prepare high-quality data for supervised, unsupervised, and deep learning models.
- Developed regression, classification, clustering, anomaly detection, and recommendation models to support customer analytics, operational optimization, business forecasting, and intelligent decision-making across enterprise applications.
- Designed and implemented transformer-based Natural Language Processing (NLP) solutions using Hugging Face Transformers and open-source LLMs (LLaMA) for text classification, text generation, sentiment analysis, and document intelligence.
- Built RESTful APIs using FastAPI and deployed production-ready machine learning models using Docker, Kubernetes, AWS SageMaker, and Apache Airflow to enable scalable AI solutions across enterprise environments.
- Conducted model evaluation using statistical analysis, hypothesis testing, experimental design, Accuracy, Precision, Recall, F1-Score, ROC-AUC, and RMSE to identify optimal models and support data-driven business decisions.
- Collaborated with client stakeholders, business analysts, product teams, and engineers to gather customer requirements, assess existing systems and workflows, identify integration opportunities, translate requirements into technical specifications, and deliver AI-driven solutions aligned with business objectives.

Environment: Python, SQL, Pandas, NumPy, Scikit-learn, TensorFlow, PyTorch, Keras, Hugging Face Transformers, FastAPI, AWS (S3, EC2, SageMaker, Glue, Lambda), Docker, Kubernetes, Apache Airflow, Tableau, Power BI, Plotly, Git, GitHub.

Education

Master of Business Administration (MBA)

Lewis University, Romeoville, IL- Jan 2024 – May 2025

Certifications

AWS Certified Machine Learning - Associate